

1 File system

1.1 Génération

Squelette de rootfs dans `workspace/nano/buildroot/system/skeleton`. Il est ensuite copié dans `buildroot/output/target` et les fichiers nécessaires y sont ensuite ajoutés.

Une fois que tous les fichiers sont ajoutés, une image `rootfs.xxx` est créé (xxx est ext4, squashfs, etc...)

1.2 1. De connaître les différents types de systèmes de fichiers ainsi que leurs applications

Pour les systèmes embarqués, il existe deux catégories de systèmes de fichiers :

1. Volatiles en RAM
2. Persistants sur des Flash (NOR et de plus en plus NAND)

Deux technologies principales sont disponible sur les Flash :

- MTD (Memory Technology Device)
- MMC/SD-Card (Multi-Media-Card / Secure Digital Card)

1.2.1 FS types

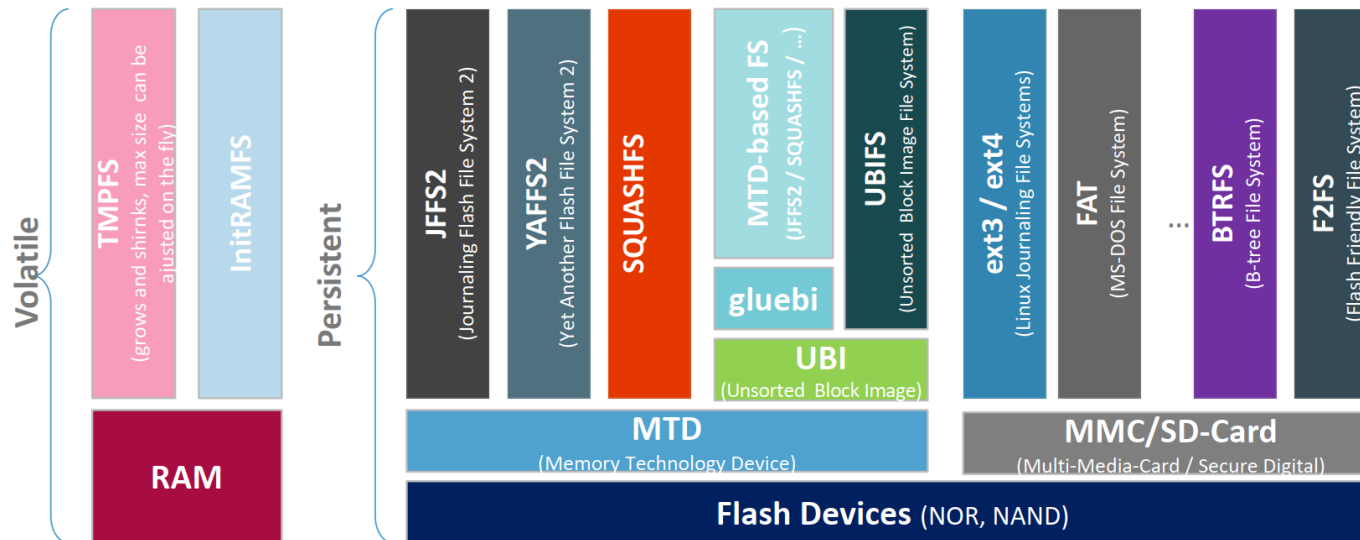


Figure 1: FS type

1.2.2 Choix d'un FS

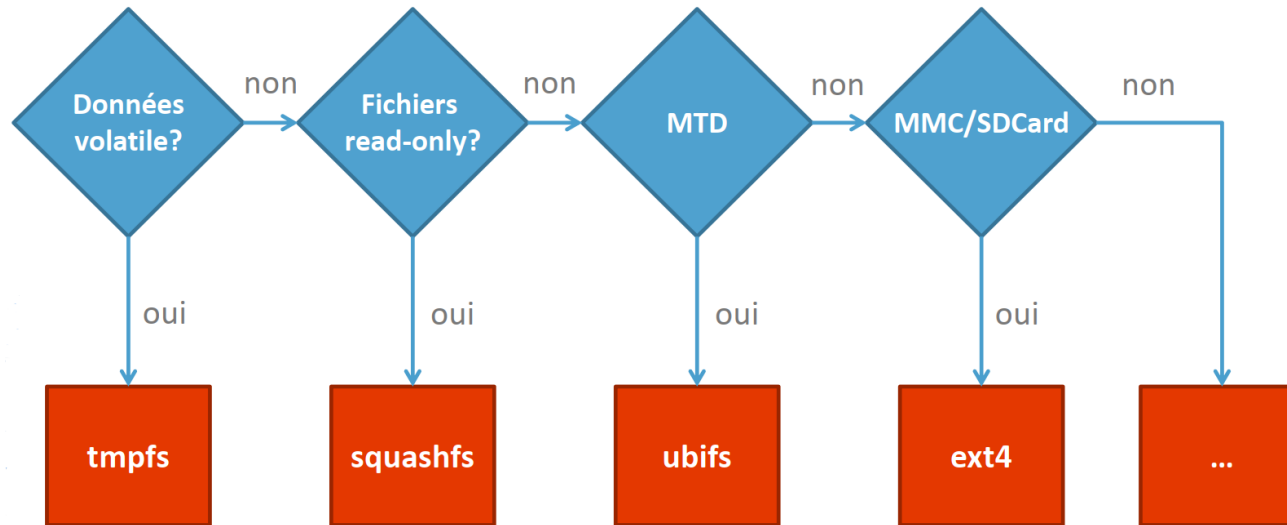


Figure 2: FS type

1.2.3 MMC technologies

MMC-eMMC-SD Card is composed by 3 elements

- MMC interface: handle communication with host
- FTL (Flash translation layer)
- Storage area: array of NAND chips

1.2.4 FTL

FTL is a small controller running a firmware. Its main purpose is to transform logical sector addressing into NAND addressing. It also handles:

- Bad block management
- Garbage collection.
- Wear levelling

1.3 2. De connaître les caractéristiques des filesystems ext2-3-4, ainsi que les commandes associées

“Filesystem considerations for embedded devices” is a good study about filesystems used on embedded systems.

This file system is very used in different Linux distribution.

- EXT filesystem was created in April 1992 and is a file system for the Linux kernel
 - Ext2 is not a journaled file system
 - Ext2 uses block mapping in order to reduce file fragmentation (it allocates several free blocks)

- After an unexpected power failure or system crash (also called an unclean system shutdown), each mounted ext2 file system on the machine must be checked for consistency with the e2fsck program
- EXT2 replaced it in 1993
 - It was merged in the 2.4.15 kernel on November 2001
 - Ext3 is compatible with ext2
 - Ext3 is a journaled file system
 - The ext3 file system prevents loss of data integrity even when an unclean system shutdown occurs
- EXT4 arrived as a stable version in the Linux kernel in 2008
 - ext4 is backward compatible with ext3 and ext2, making it possible to mount ext3 and ext2 as ext4
 - Ext4 is included in the kernel 2.6.28
 - Ext4 supports Large file system:
 - * Volume max: 2^{60} bytes
 - * File max: 2^{40} bytes
 - Ext4 uses extents (as opposed to the traditional block mapping scheme used by ext2 and ext3), which improves performance when using large files and reduces metadata overhead for large files

1.3.1 Ext4 commands

```
# Create a partition (rootfs), start 64MB, length 256MB
sudo parted /dev/sdb mkpart primary ext4 131072s 655359s
# Format the partition with the volume label = rootfs
sudo mkfs.ext4 /dev/sdb1 -L rootfs
# Modify (on the fly) the ext4 configuration
sudo tune2fs <options> /dev/sdb1
# check the ext4 configuration
mount
sudo tune2fs -l /dev/sdb1
sudo dumpe2fs /dev/sdb1
# mount an ext4 file system
mount -t ext4 /dev/sdb1 /mnt/test // with default options
mount -t ext4 -o defaults,noatime,discard,nodiratime,data=writeback,acl,user_xattr
/dev/sdb1 /mnt/test
```

1.3.2 Ext4 mount options and MMC/SD-Card

- filesystem options can be activated with the mount command (or to the /etc/fstab file)
- These options can be modified with tune2fs command
- Journaling: the journaling guarantees the data consistency, but it reduces the file system performances
- MMC/SD-Card constraints: In order to improve the longevity of MMC/SDCard, it is necessary to reduce the unnecessary writes
- Mount options to reduce the unnecessary writes (man mount) :
 - noatime: Do not update inode access times on this filesystem (e.g., for faster access on the news spool to speed up news servers)
 - nodiratime: Do not update directory inode access times on this filesystem
 - relatime: this option can replace the noatime and nodiratime if an application needs the access time information (like mutt)

Mount options for the journaling (man ext4):

- Data=journal: All data is committed into the journal prior to being written into the main filesystem (It is the safest option in terms of data integrity and reliability, though maybe not so much for performance)
- Data=ordered: This is the default mode. All data is forced directly out to the main file system before the metadata being committed to the journal
- Data=writeback: Data ordering is not preserved - data may be written into the main filesystem after its metadata has been committed to the journal. It guarantees internal filesystem integrity, however it can allow old data to appear in files after a crash and journal recovery.
- Discard: Use discard requests to inform the storage that a given range of blocks is no longer in use. A MMC/SD-Card can use this information to free up space internally, using the free blocks for wear-levelling.
- acl: Support POSIX Access Control Lists
- user_xattr: Support "user." extended attributes
- default : rw, suid, dev, exec, auto, nouser, and async
 - rw : read-write
 - suid : Allow set-user-identifier or set-group-identifier bits
 - dev : Interpret character or block special devices on the filesystem
 - exec : Permit execution of binaries
 - auto : Can be mounted with the -a option (mount -a)
 - nouser : Forbid an ordinary (i.e., non-root) user to mount the filesystem
 - async : All I/O to the filesystem should be done asynchronously

1.3.3 /etc/fstab file

File **/etc/fstab** contains descriptive information about the filesystems the system can mount

- `{file system}_i` : block special device or remote filesystem to be mounted
- `{mount pt}_i` : mount point for the filesystem
- `{type}_i` : the filesystem type
- `{options}_i` : mount options associated with the filesystem
- `{dump}_i` : used by the dump (backup filesystem) command to determine which filesystems need to be dumped (0 -i no backup)
- `{pass}_i` : used by the fsck (8) program to determine the order in which filesystem checks are done at reboot time. The root filesystem should be specified with 1, and other filesystems should have a 2. if `{pass}_i` is not present or equal 0 -i fsck will assume that the filesystem is not checked.
- Field options: It contains at least the type of mount plus any additional options appropriate to the filesystem type.

Common for all types of file system are the options (man mount) :

- auto : Can be mounted with the -a option (mount -a)
- defaults : Use default options: rw, suid, dev, exec, auto, nouser, and async
- nosuid : Do not allow set-user-identifier or set-group-identifier bits to take effect
- noexec : Do not allow direct execution of any binaries on the mounted file system
- nodev : Do not interpret character or block special devices on the file system

1.4 3. D'expliquer les différents "files systems" utilisés dans les systèmes embarqués (ext2-3-4, BTRFS, F2FS, NILFS2, XFS, ZFS, ...)

1.4.1 BTRFS (B-Tree filesystem)

- BTRFS is a "new" file system compared to EXT. It is originally created by Oracle in 2007, it is a B-Tree filesystem
- It is considered stable since 2014
- Since 2015 BTRFS is the default rootfs for openSUSE
- BTRFS inspires from both Reiserfs and ZFS
- Theodore Ts'o (ext3-ext4 main developer) said that BTRFS has a better direction than ext4 because "it offers improvements in scalability, reliability, and ease of management"

1.4.2 F2FS (Flash-Friendly File System)

It is a log filesystem. It can be tuned using many parameters to allow best handling on different supports.

F2FS features :

- Atomic operations
- Defragmentation
- TRIM support (reporting free blocks for reuse)

1.4.3 NILFS2 (New Implementation of a Log-structured File System)

- Developed by Nippon Telegraph and Telephone Corporation
- NILFS2 Merged in Linux kernel version 2.6.30
- NILFS2 is a log filesystem
- CoW for checkpoints and snapshots
- Userspace garbage collector

1.4.4 XFS (Flash-Friendly File System)

XFS was developed by SGI in 1993.

- Added to Linux kernel in 2001
- On disk format updated in Linux version 3.10
- XFS is a journaling filesystem
- Supports huge filesystems
- Designed for scalability
- Does not seem to be handling power loss (standby state) well

1.4.5 ZFS (Zettabyte (10²¹)File System)

ZFS is a combined file system and logical volume manager designed by Sun Microsystems.

- ZFS is a B-Tree file system
- Provides strong data integrity
- Supports huge filesystems
- Not intended for embedded systems (requires RAM)
- License not compatible with Linux

1.4.6 Conclusion

Performances:

- EXT4 is currently the best solution for embedded systems using MMC
- F2FS and NILFS2 show impressive write performances

Features:

- BTRFS is a next generation filesystem
- NILFS2 provides simpler but similar features

Scalability:

- EXT4 clearly doesn't scale as well as BTRFS and F2FS




Figure 3: FS Comparaison

1.5 4. Expliquer les files system de type Journal, B-Tree/CoW, log filesystem

1.5.1 Journalized filesystem

A journalized filesystem keeps track of every modification in a journal in a dedicated area

- EXT3, EXT4, XFS, Reiser4
- Journal allows to restore a corrupted filesystem
- Modifications are first recorded in the journal
- Modifications are applied on the disk
- If a corruption occurs: The File System will either keep or drop the modifications
 - Journal is consistent : we replay the journal at mount time
 - Journal is not consistent : we drop the modifications

1.5.2 B-Tree filesystem

- ZFS, BTRFS, NILFS2

- B+ tree is a data structure that generalized binary trees
- CoW (Copy on Write) is used to ensure no corruption occurs at runtime :
 - The original storage is never modified. When a write request is made, data is written to a new storage area
 - Original storage is preserved until modifications are committed
 - If an interruption occurs during writing the new storage area, original storage can be used

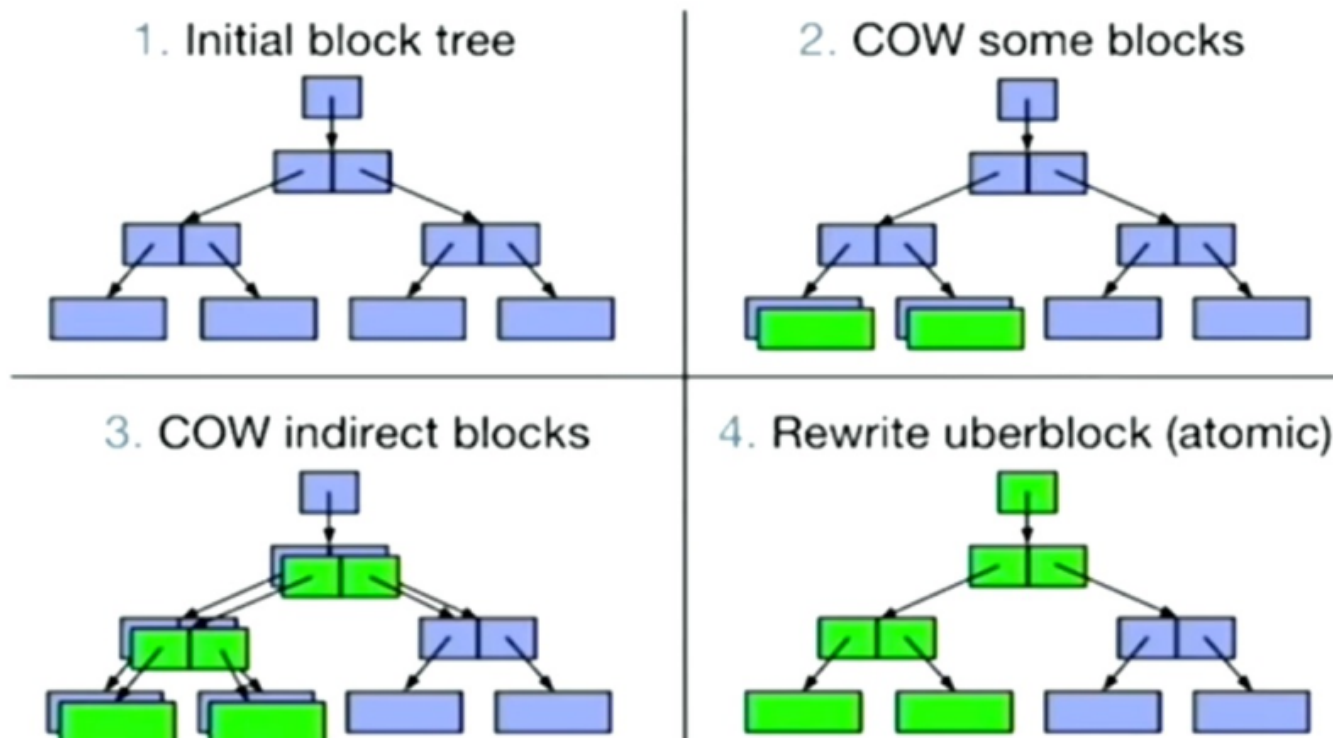


Figure 4: B-tree type FS execution

1.5.3 Log filesystem

Log-structured filesystems use the storage medium as circular buffer and new blocks are always written to the end.

- F2FS, NILFS2, JFFS2, UBIFS
- Log-structured filesystems are often used for flash media since they will naturally perform wear-levelling
- The log-structured approach is a specific form of copy-on-write behavior

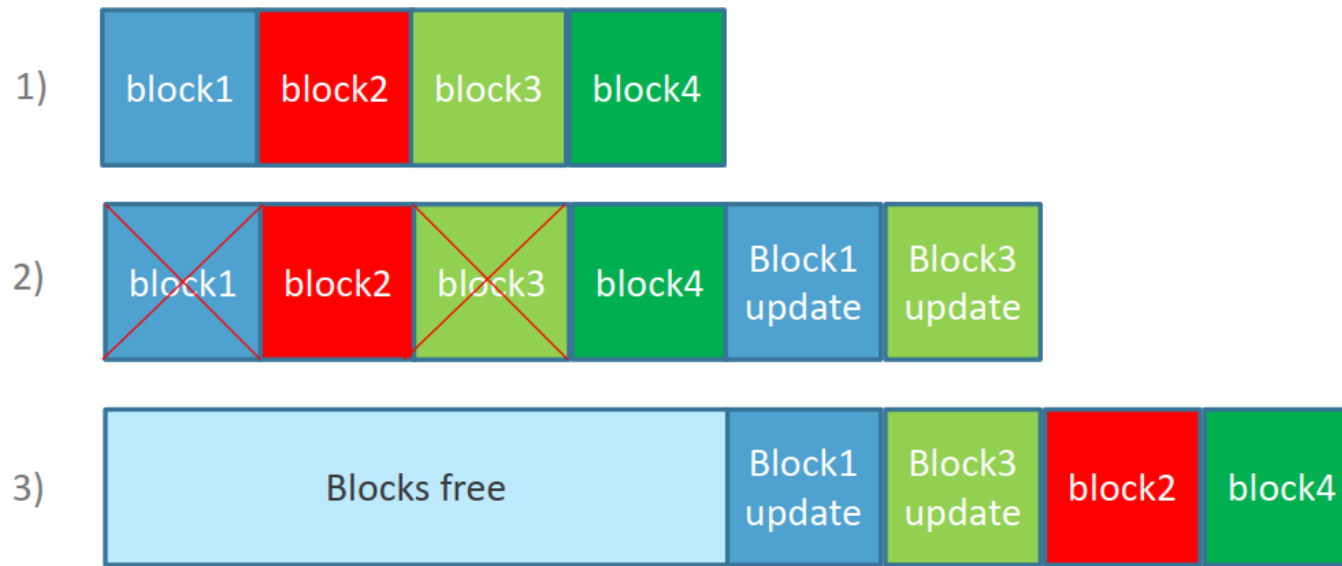


Figure 5: log_fs type FS execution

1. Initial state
2. Block 1-3 are updated, old blocks 1-3 are not used
3. Garbage copies block2 and 4, and frees old block1-2-3-4

1.6 5. De connaître les caractéristiques du filesystem Squashfs, ainsi que les commandes associées

beginitemize

Squashfs is a compressed read-only filesystem for Linux

Squashfs versions

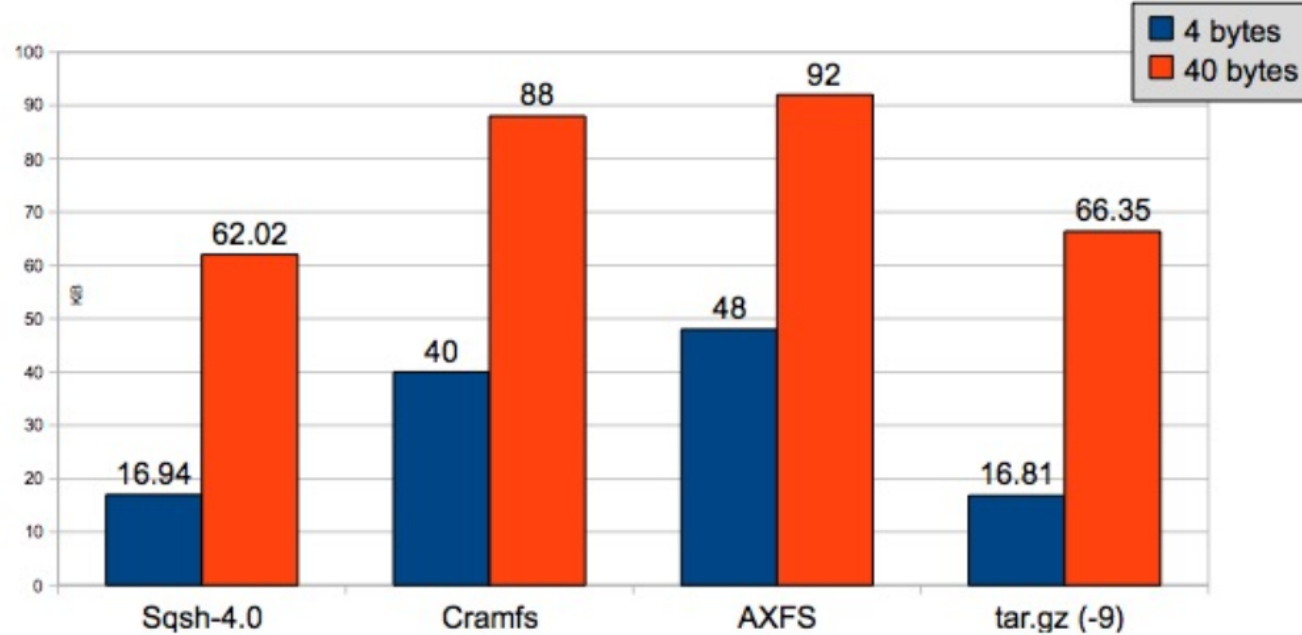
- Squashfs 2.0 and squashfs 2.1: 2004, kernel 2.2
- Squashfs 3.0: 2006, kernel 2.6.12
- Squashfs 4.2: 2011, kernel 2.6.29

It uses gzip, lzma, lzo, lz4 and xz compression to compress files, inodes and directories

SquashFS 4.0 supports 64 bit filesystems and files (larger than 4GB), full uid/gid information, hard links and timestamps

Squashfs is intended for general read-only filesystem use, for archival use, and in embedded systems with small processors where low overhead is needed

1200 small files size with different compression techniques



1.6.1 Create and use squashed file systems

1. Create the squashed file system dir.sqsh for the regular directory /data/config/ :

```
bash# mksquashfs /data/config/ /data/dir.sqsh
```



2. The mount command is used with a loopback device in order to read the squashed file system dir.sqsh

```
bash# mkdir /mnt/dir
bash# mount -o loop -t squashfs /data/dir.sqsh /mnt/dir
bash# ls /mnt/dir
```

3. It is possible to copy the dir.sqsh to an unmounted partition (e.g. /dev/sdb2) with the dd command and next to mount the partition as squashfs file system

```
bash# umount /dev/sdb2
bash# dd if=dir.sqsh of=/dev/sdb2
bash# mount /dev/sdb2 /mnt/dir -t squashfs
bash# ls /mnt/dir
```

1.7 6. De connaître les caractéristiques du filesystem tmpfs, ainsi que les commandes associées

- Tmpfs is a file system which keeps all files in virtual memory
- Everything in tmpfs is temporary in the sense that no files will be created on your hard drive. If you unmount a tmpfs instance, everything stored therein is lost.
- tmpfs puts everything into the kernel internal caches and grows and shrinks to accommodate the files it contains and is able to swap unneeded pages out to swap space. It has maximum size limits which can be adjusted on the fly via 'mount -o remount ...'
- If you compare it to ramfs you gain swapping and limit checking. Another similar thing is the RAM disk (/dev/ram*), which simulates a fixed size hard disk in physical RAM, where you have to create an ordinary filesystem on top. Ramdisks cannot swap and you do not have the possibility to resize them

1.7.1 Devtmpfs

- devtmpfs is a file system with automatically populates nodes files (/dev/...) known by the kernel.
- This means, it is not necessary to have udev running nor to create a static /dev layout with additional, unneeded and not present device nodes.
- Instead the kernel populates the appropriate information based on the known devices.
- The kernel executes this command : mount -n -t devtmpfs devtmpfs /dev
- /dev is automatically populated by the kernel with its known devices

```
# ls /dev
autofs ptyrf tty47
btrfs-control random tty48
bus rtc0 tty49
console shm tty5
cpu_dma_latency snapshot tty50
```

1.8 7. De connaître les caractéristiques du filesystem LUKS, ainsi que les commandes associées

1.9 8. Savoir expliquer la gestion des clés de LUKS 42

1.10 9. De connaître les caractéristiques du filesystem InitramFS, ainsi que les commandes associées

1.11 10. De savoir créer un initramFS